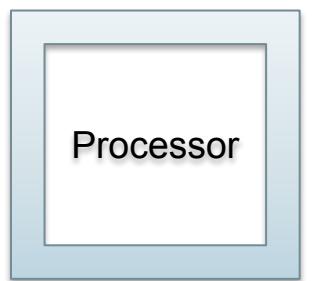
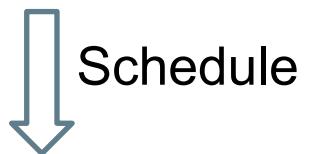
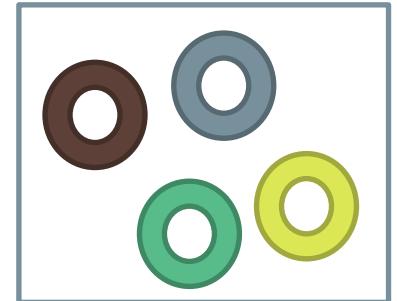


OPERATING SYSTEMS
CPU Scheduling

Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- Continuous Cycle :
 - one process has to wait (I/O)
 - Operating system takes the CPU away
 - Give CPU to another process
 - This pattern continues
- CPU-I/O Burst Cycle : Process execution consists of a *cycle* of CPU execution and I/O wait



CPU Scheduler

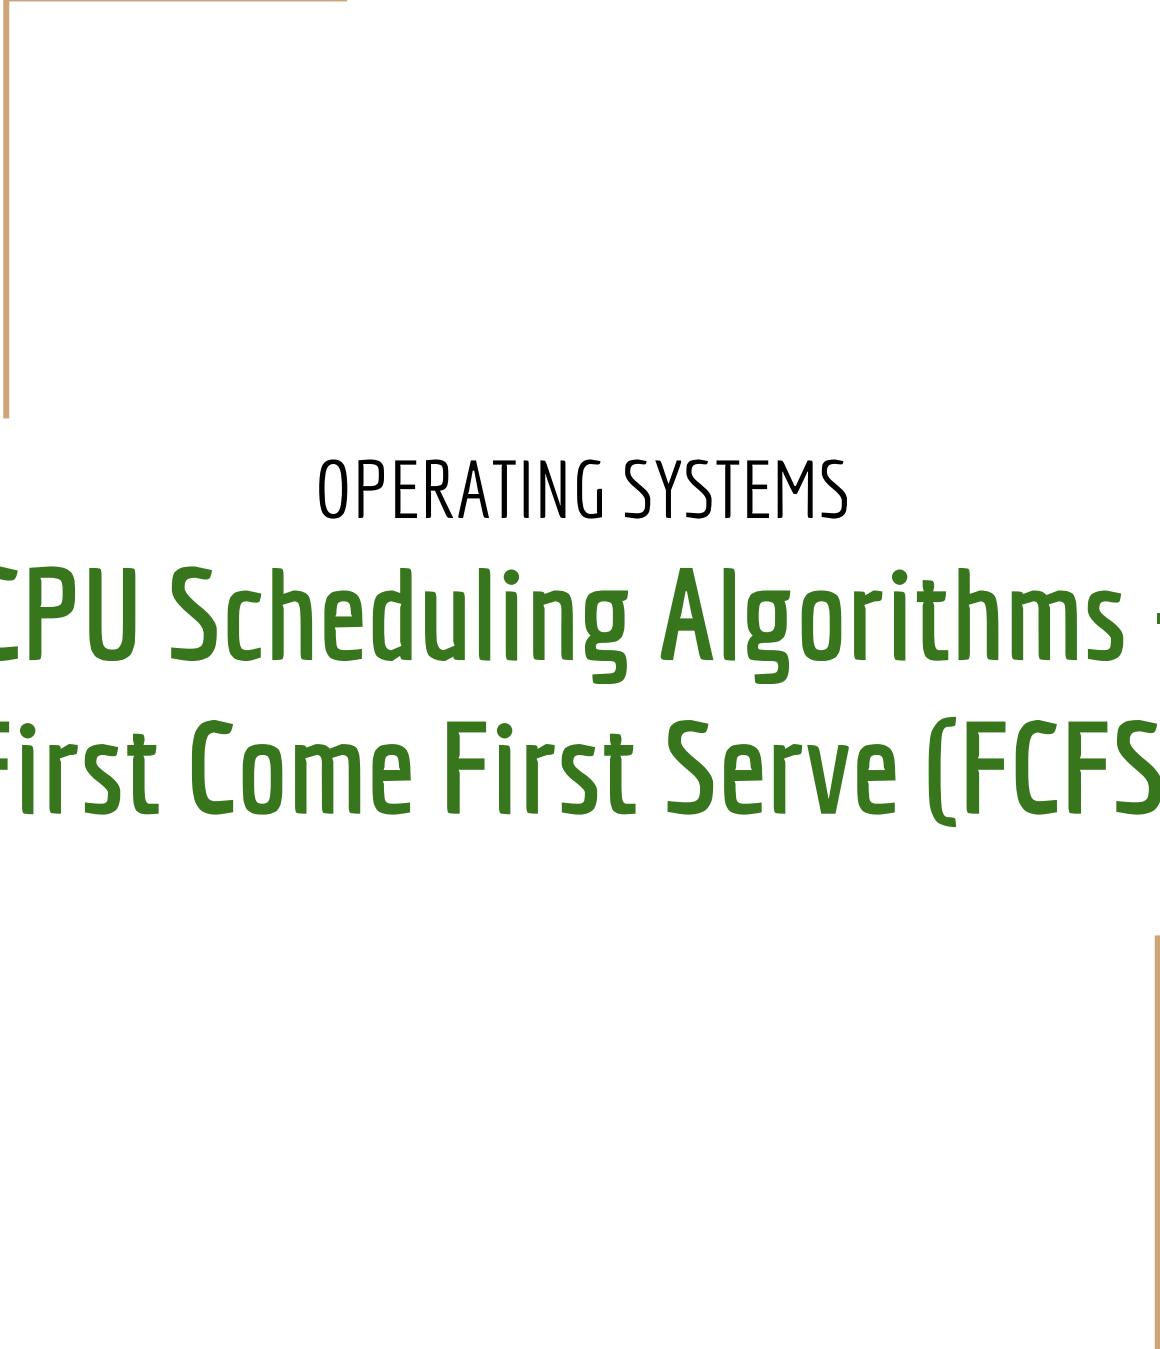
- **Selects from among the processes in ready queue, and allocates the CPU to one of them**
 - FIFO queue
 - Priority queue
 - Tree
 - Unordered linked-list
- **CPU scheduling decisions may take place when a process:**
 1. Switches from running to waiting state (I/O request)
 2. Switches from running to ready state (e.g. when interrupt occurs)
 3. Switches from waiting to ready (e.g. at completion of I/O)
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities

Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time**
 - amount of time to execute a particular process
 - the interval from the time of submission of a process to the time of the completion.
 - sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, doing I/O
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time



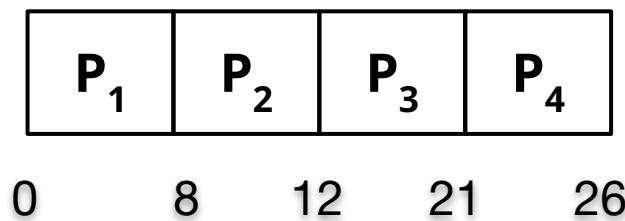
OPERATING SYSTEMS

CPU Scheduling Algorithms - First Come First Serve (FCFS)

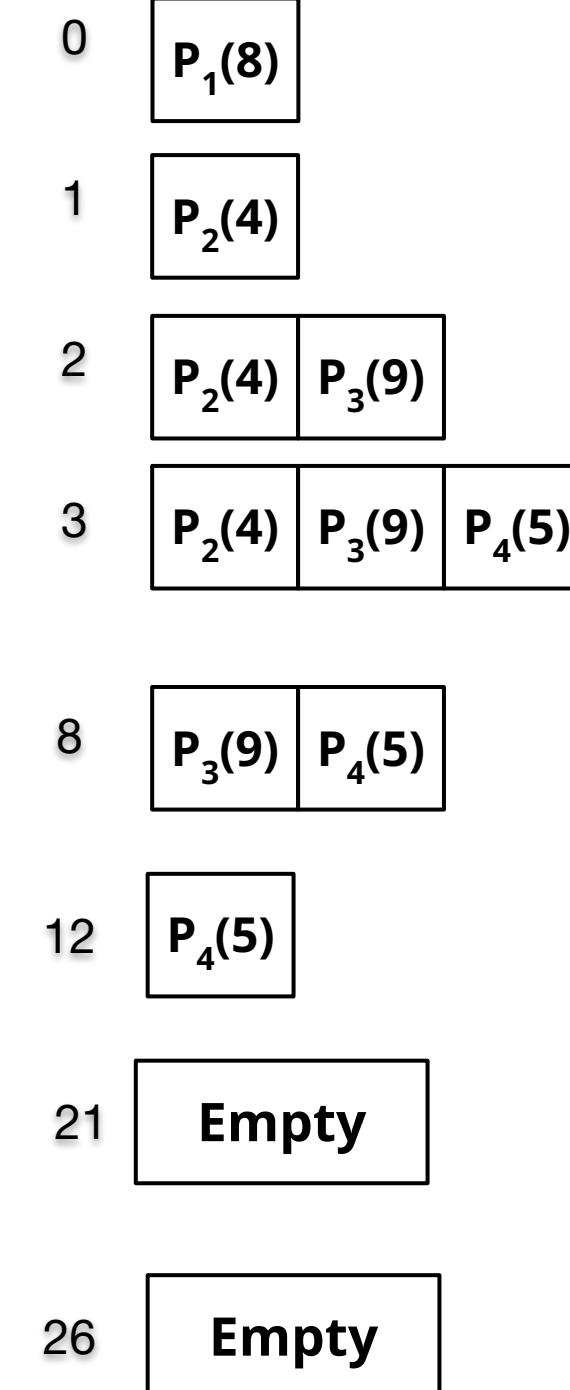
First-Come, First-Served (FCFS) Scheduling

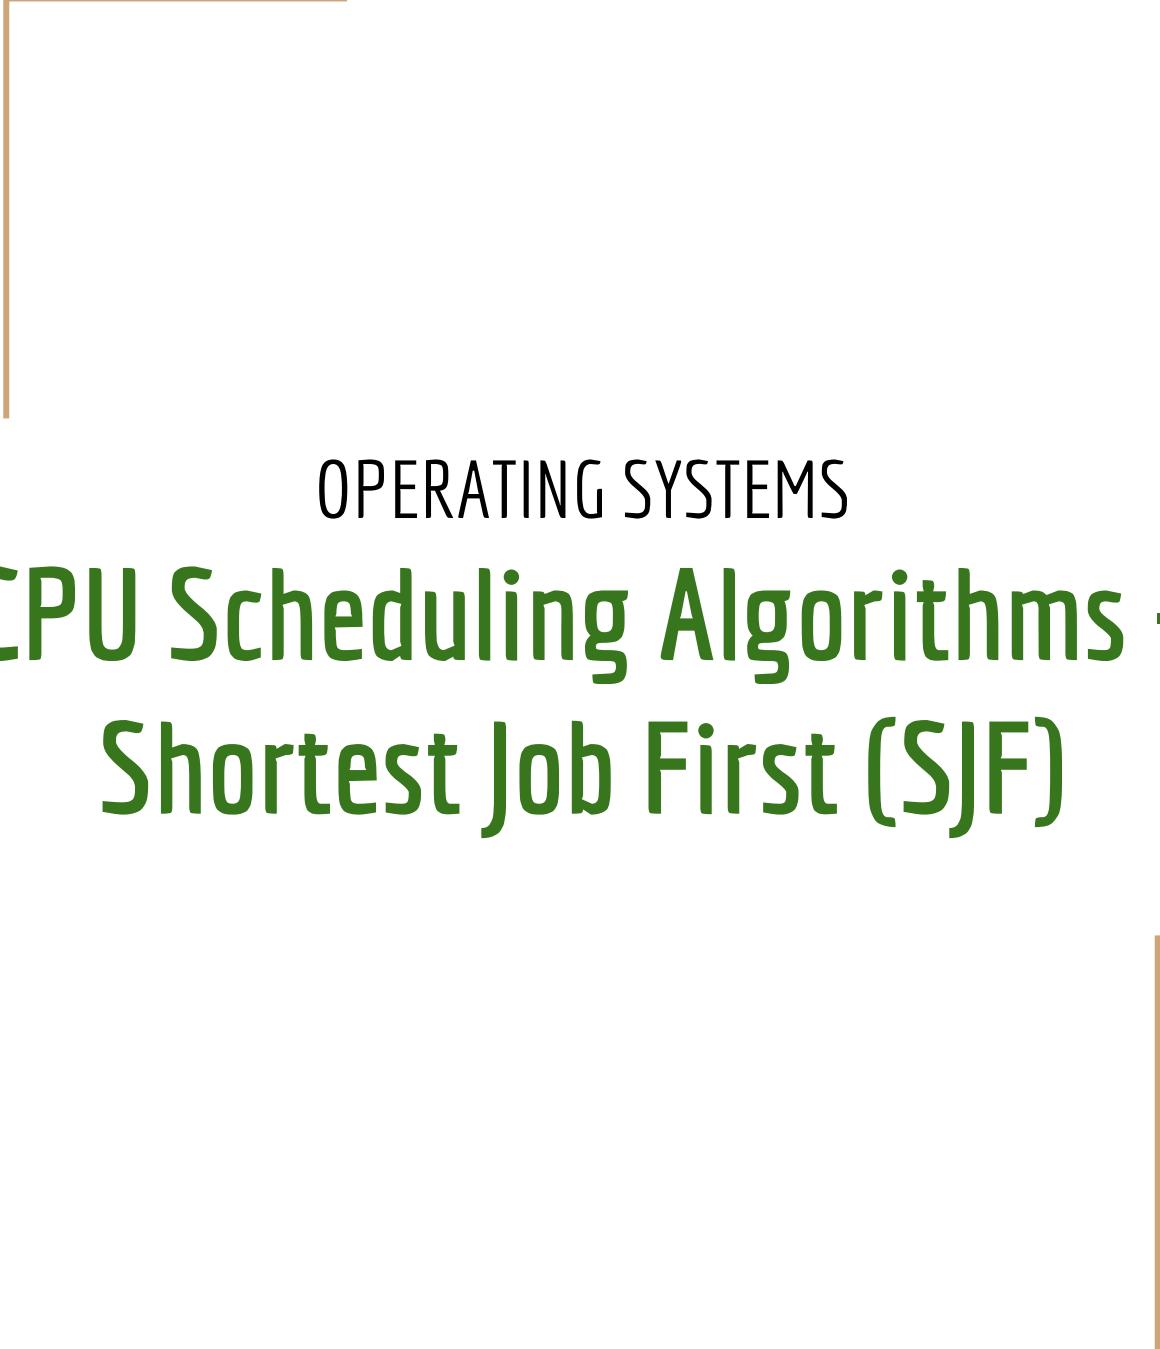
Process	Arrival Time	Burst Time
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

The Gantt Chart for the schedule is:



Ready Queue





OPERATING SYSTEMS

CPU Scheduling Algorithms - Shortest Job First (SJF)

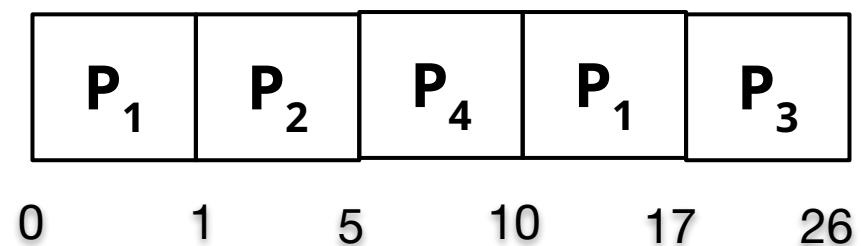
Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - **Non-preemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst
 - **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes

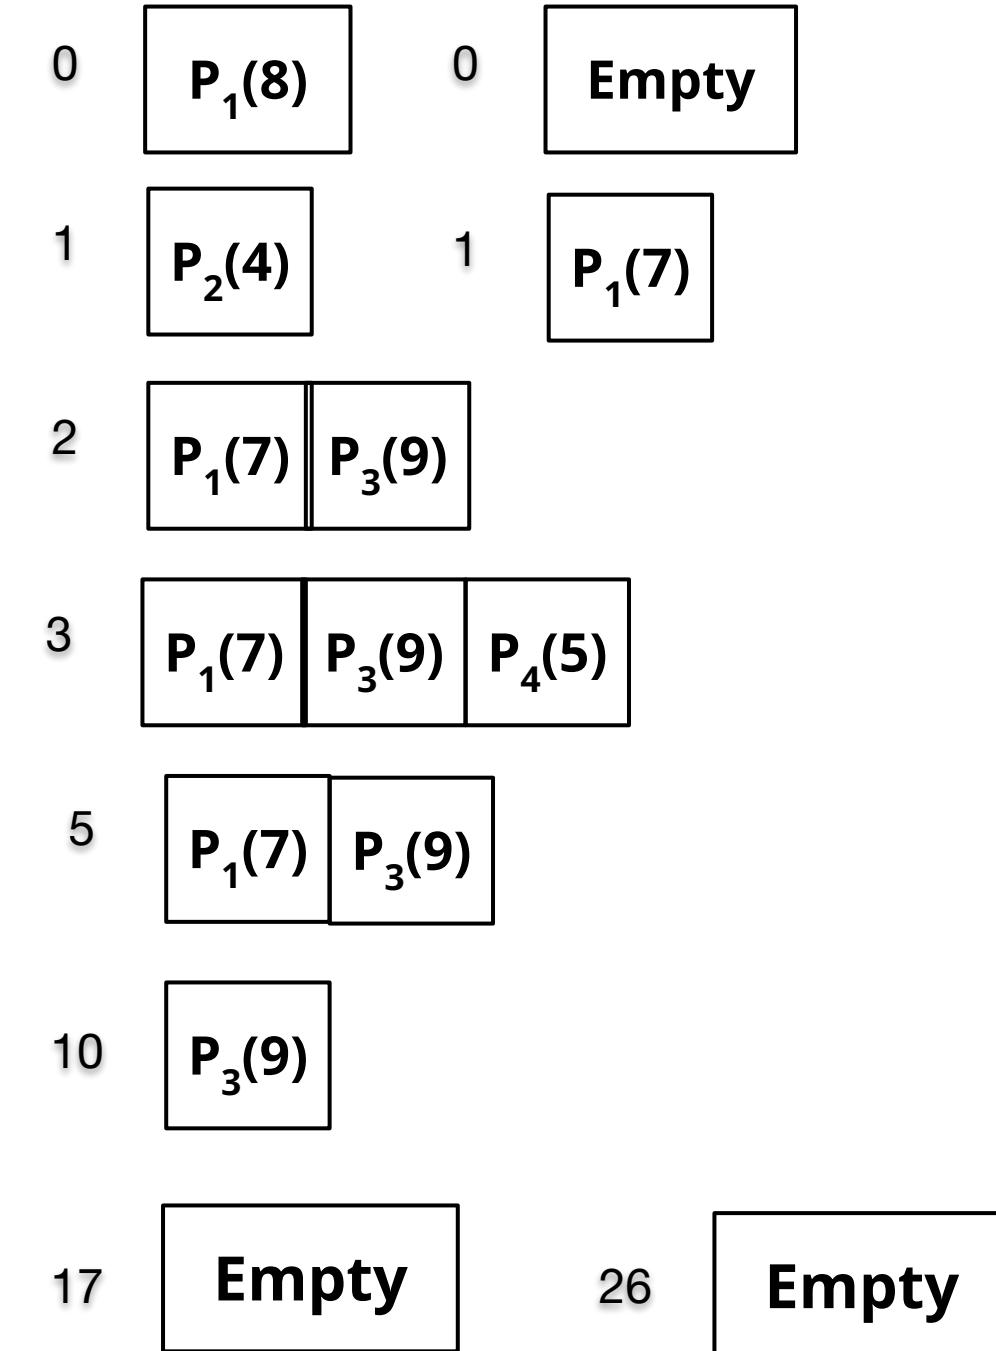
Example of Shortest-remaining-time-first

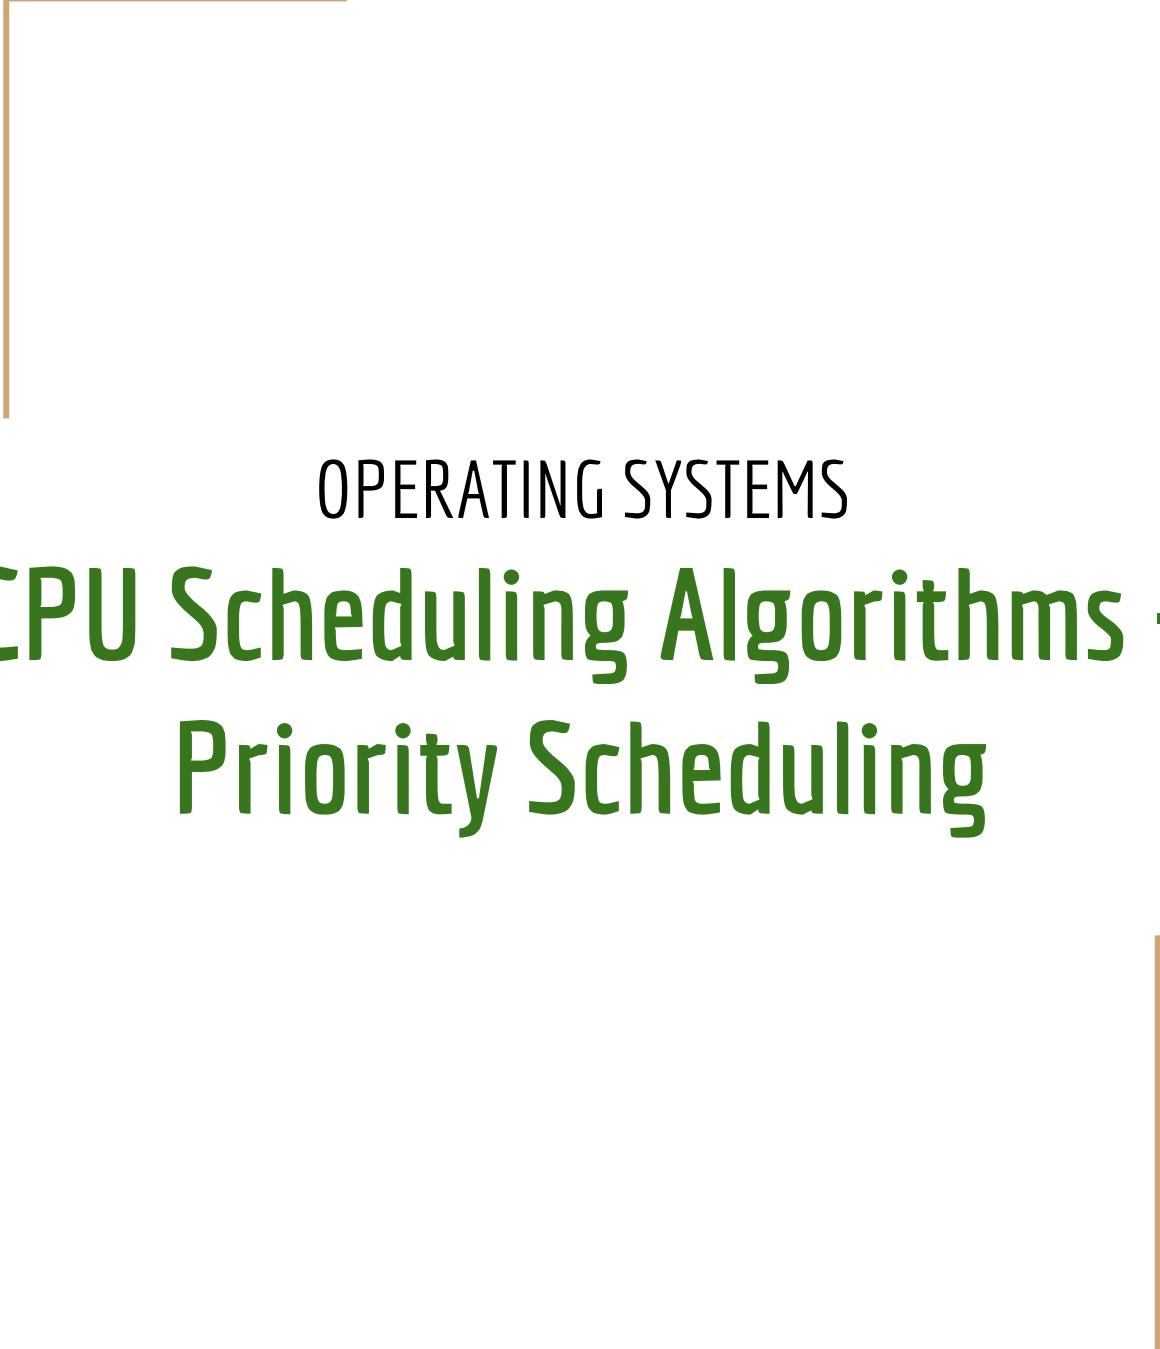
Process	Arrival Time	Burst Time
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

Gantt Chart:



Ready Queue





OPERATING SYSTEMS

CPU Scheduling Algorithms - Priority Scheduling

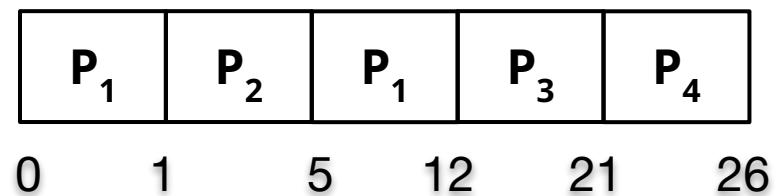
Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Priority can be defined either internally or externally.
 - Factors for internal priority assignment:
 - Time limit, memory requirements, the number of open files etc.
 - Factors for external priority assignment:
 - Importance of the process, the type and amount of funds being paid for computer use, department sponsoring works etc.

Example of Priority Scheduling (Preemptive)

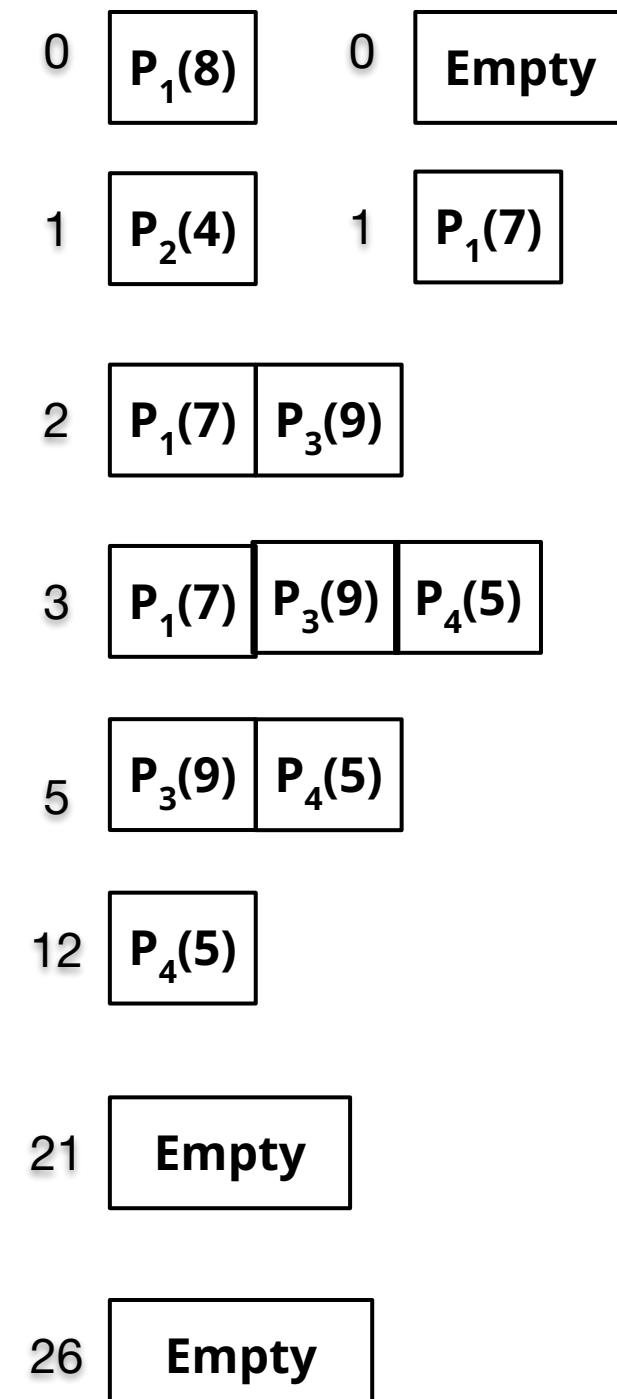
Process	Arrival Time	Burst Time	Priority
P ₁	0	8	3
P ₂	1	4	1
P ₃	2	9	4
P ₄	3	5	5

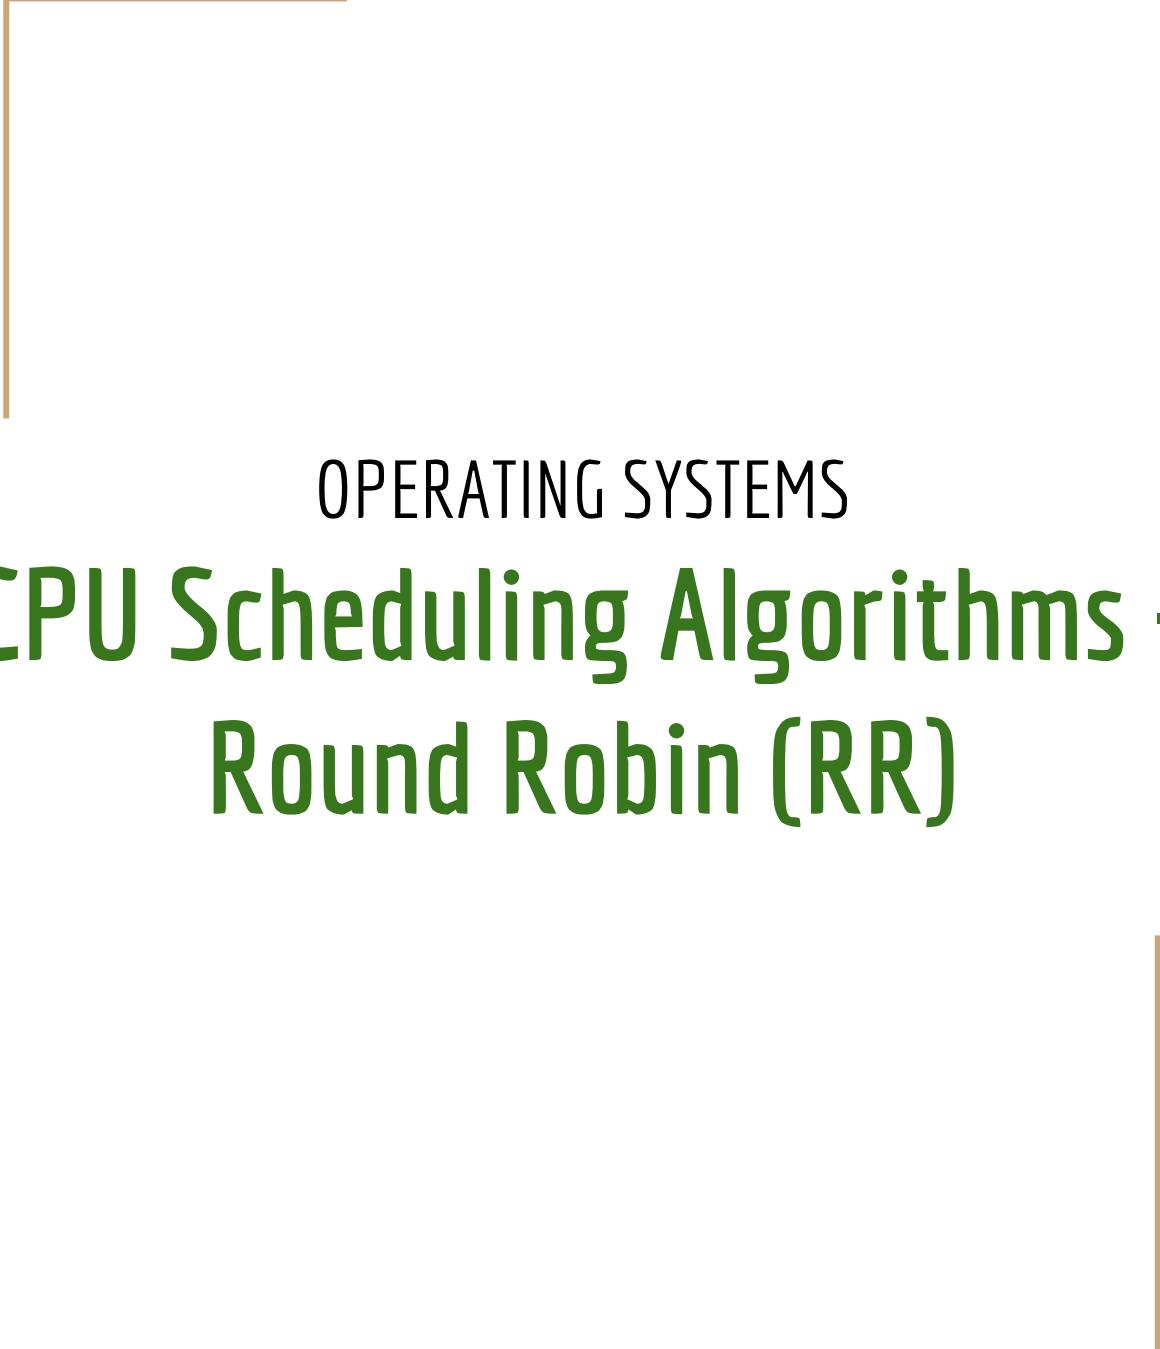
Priority scheduling Gantt Chart



- Problem ≡ **Starvation** – low priority processes may never execute
- Solution ≡ **Aging** – as time progresses increase the priority of the process

Ready Queue





OPERATING SYSTEMS

CPU Scheduling Algorithms -

Round Robin (RR)

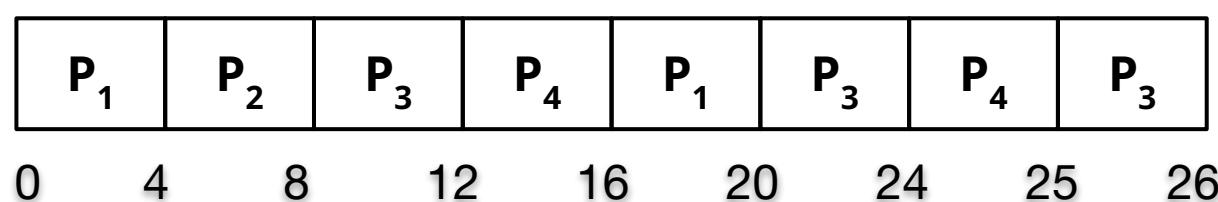
Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high

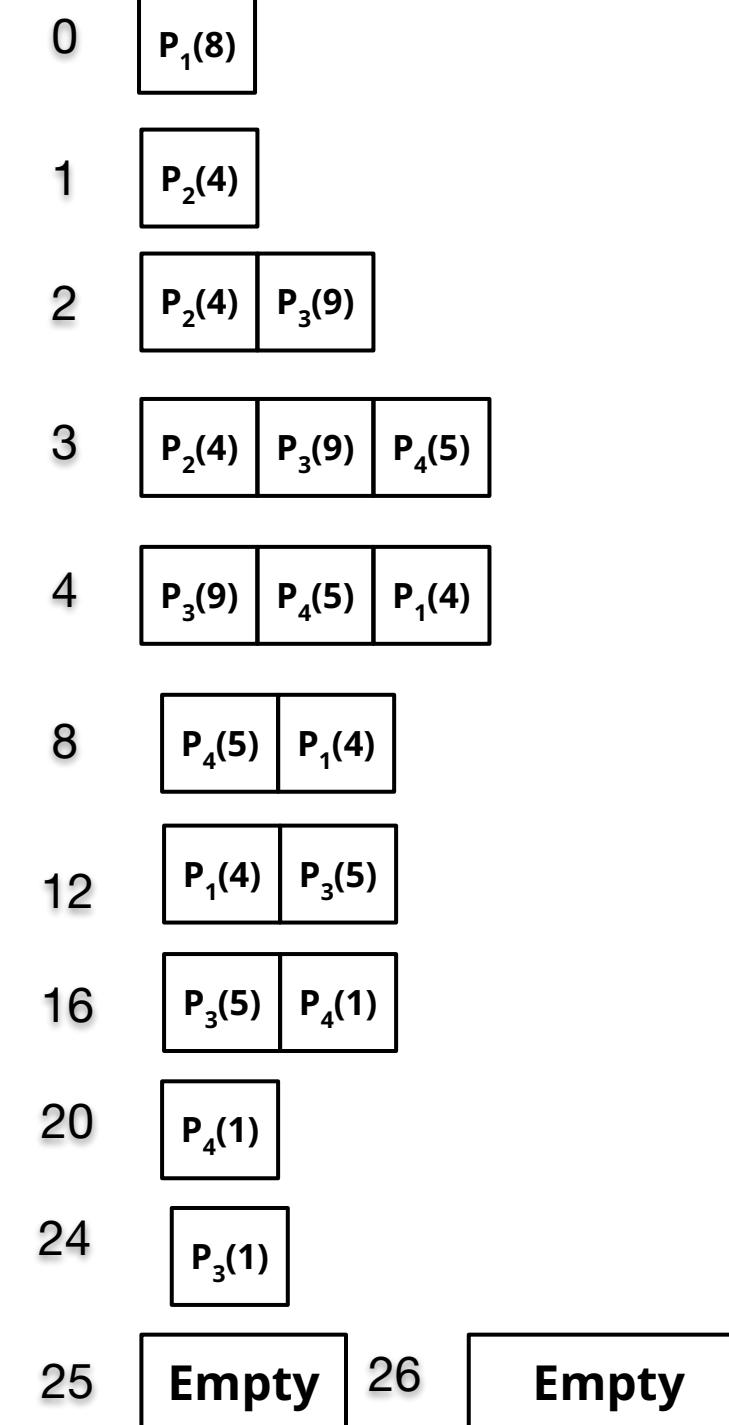
Example of RR with Time Quantum = 4

Process	Arrival Time	Burst Time
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

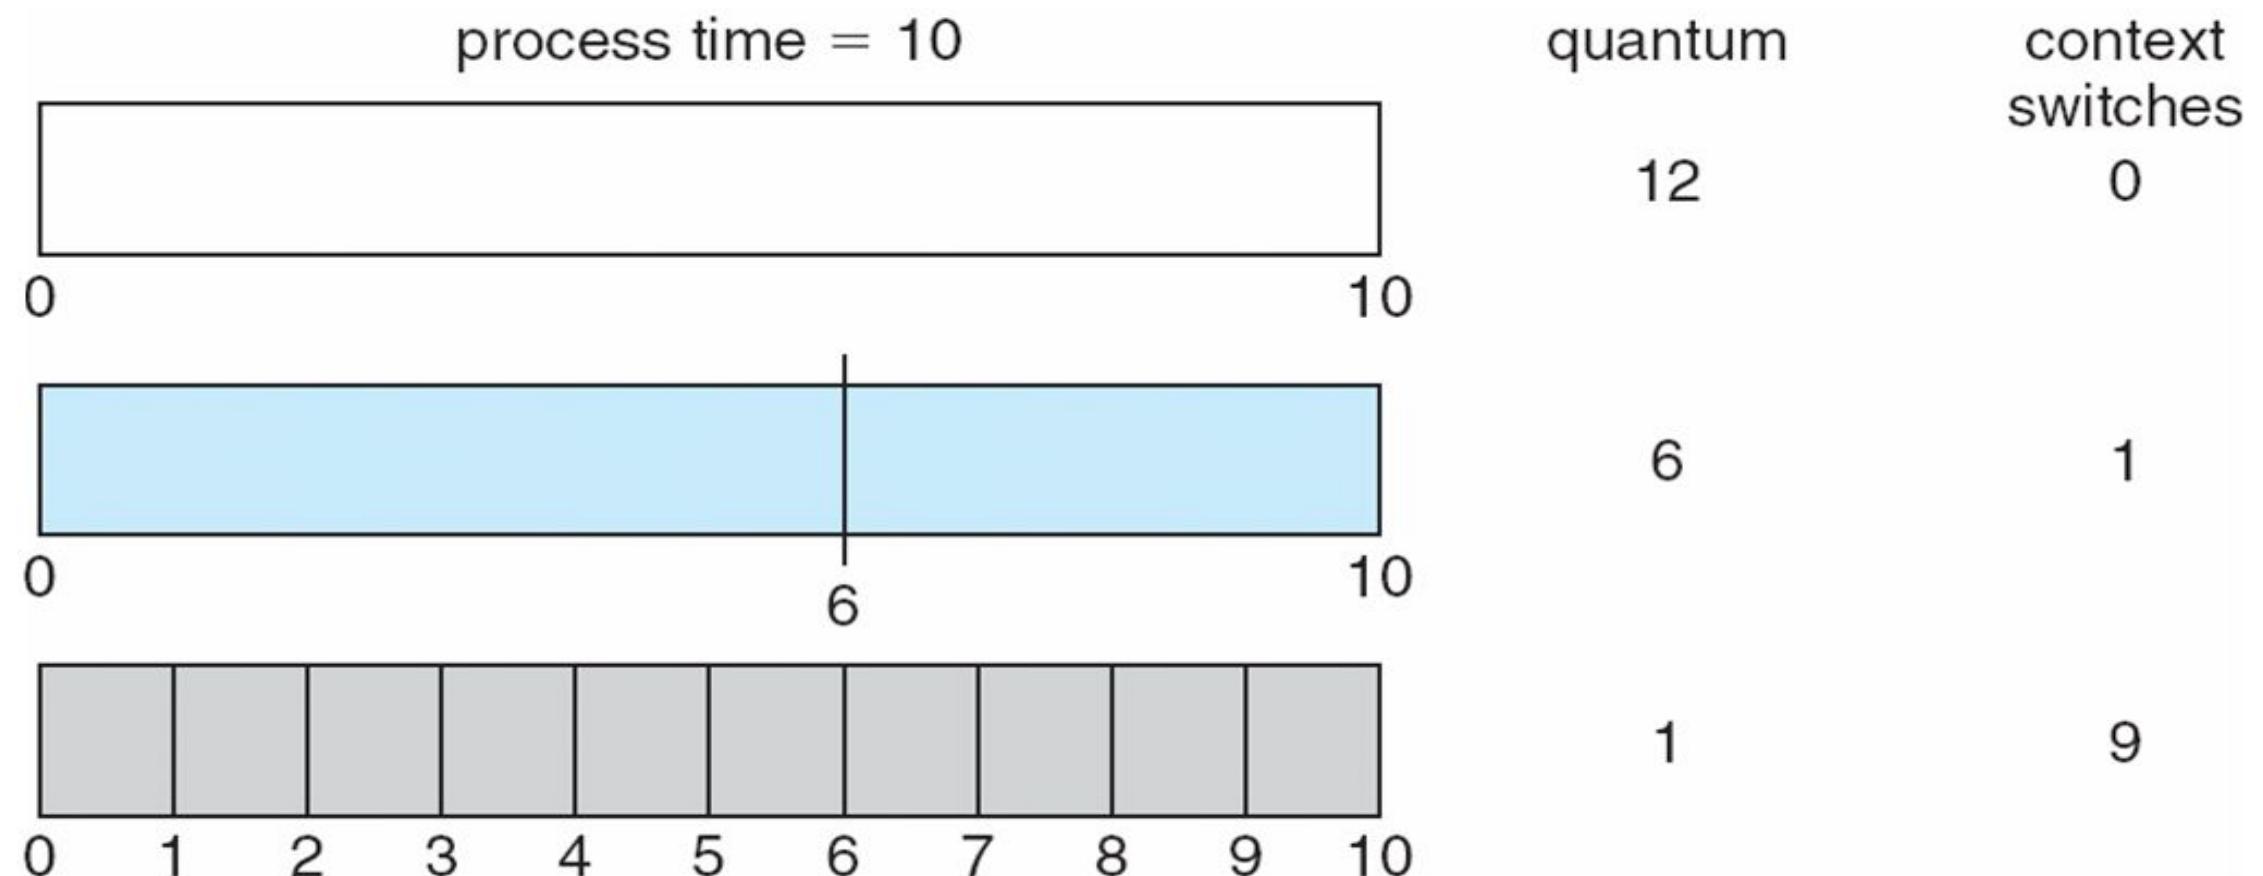
The Gantt chart is:



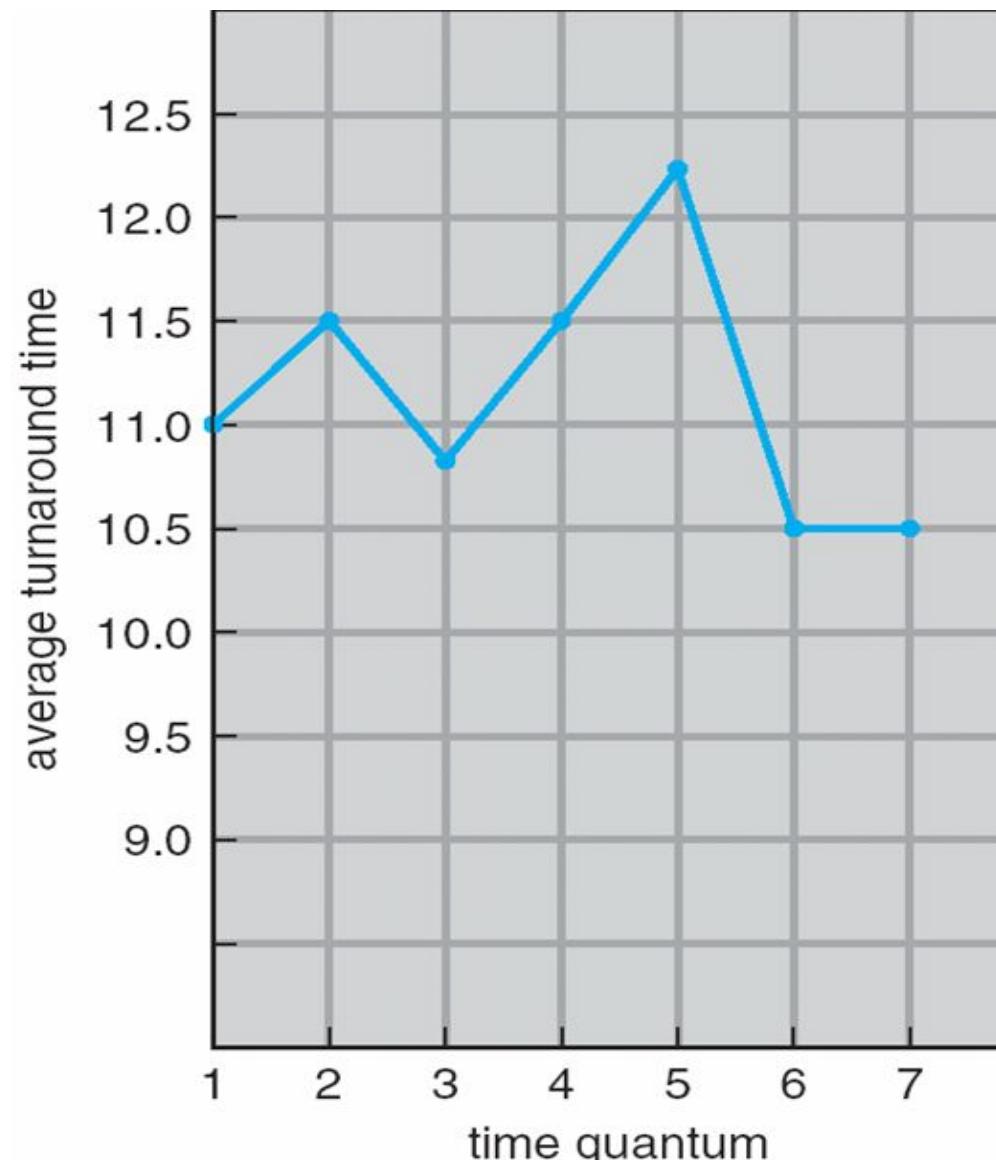
Ready Queue:



Time Quantum and Context Switch Time

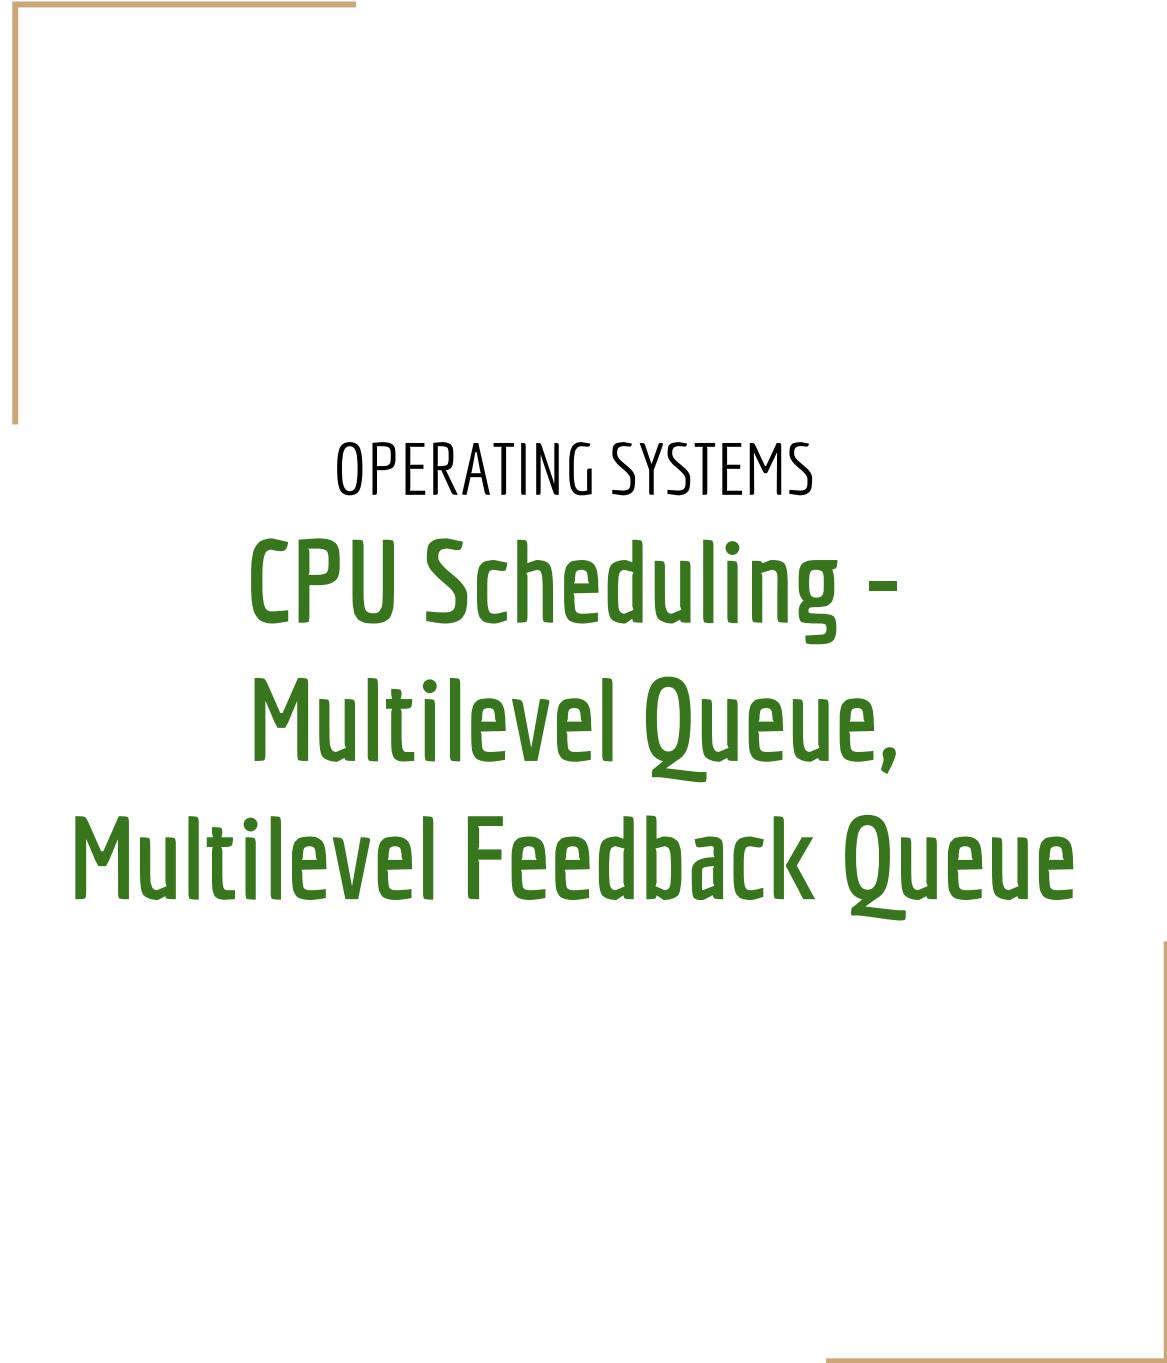


Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts should
be shorter than quantum



OPERATING SYSTEMS

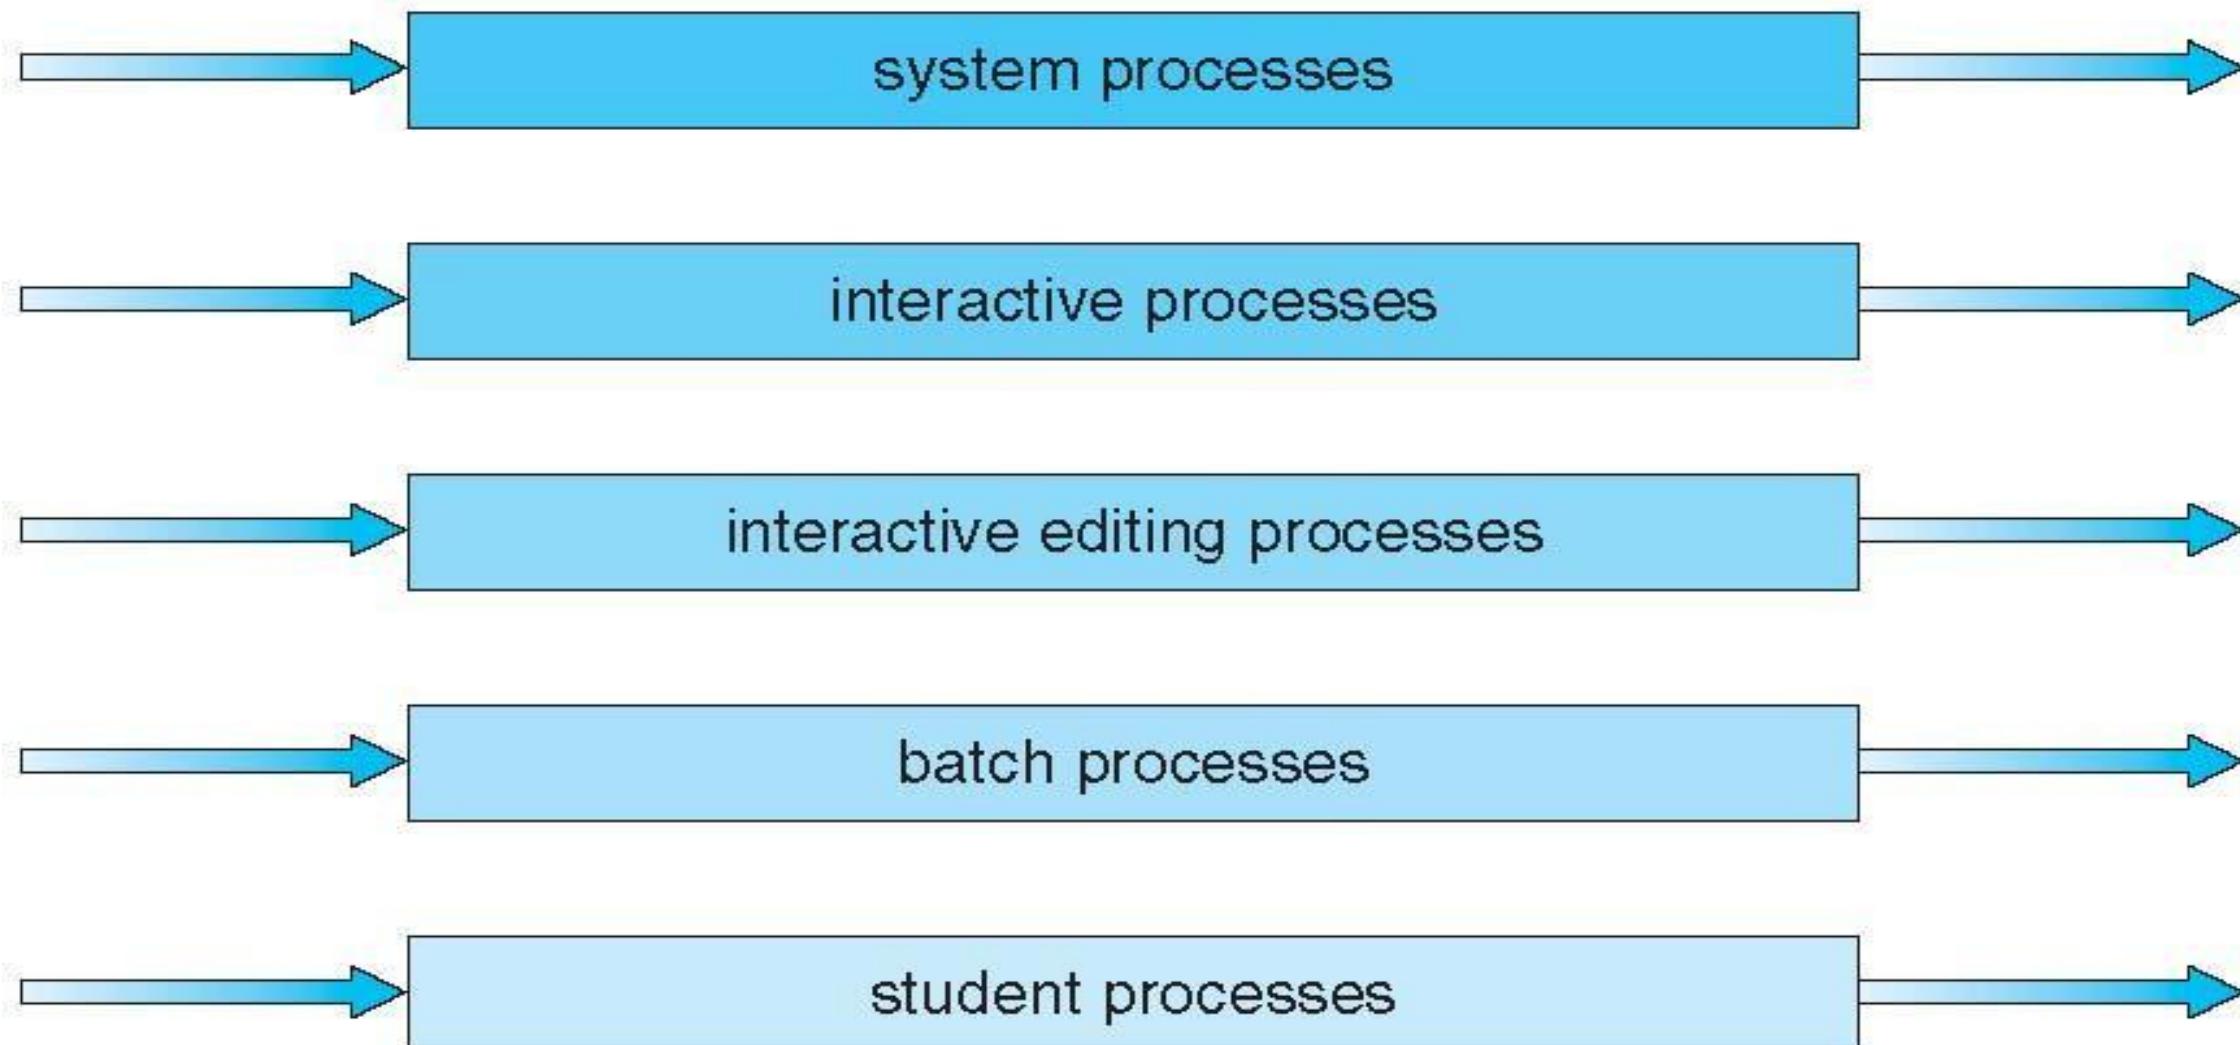
CPU Scheduling - Multilevel Queue, Multilevel Feedback Queue

Multilevel Queue

- Another class of scheduling algorithm needs- in which processes are classified into different groups, e.g.:
 - foreground (interactive) processes
 - background (batch) processes
- They have different response time requirements-so different scheduling needs.
- Foreground processes may have priority over background processes.
- A multilevel queue-scheduling algorithm partitions the ready queue into several separate queues-we can see it in the figure of next slide:-
- Each queue has its own scheduling algorithm:
 - Foreground queue scheduled by – RR algorithm
 - Background queue scheduled by – FCFS algorithm
- Scheduling must be done between the queues:
 - Fixed priority preemptive scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., foreground queue can be given 80% of the CPU time for RR-scheduling among its processes, while 20% to background in FCFS manner.

Multilevel Queue Scheduling

highest priority



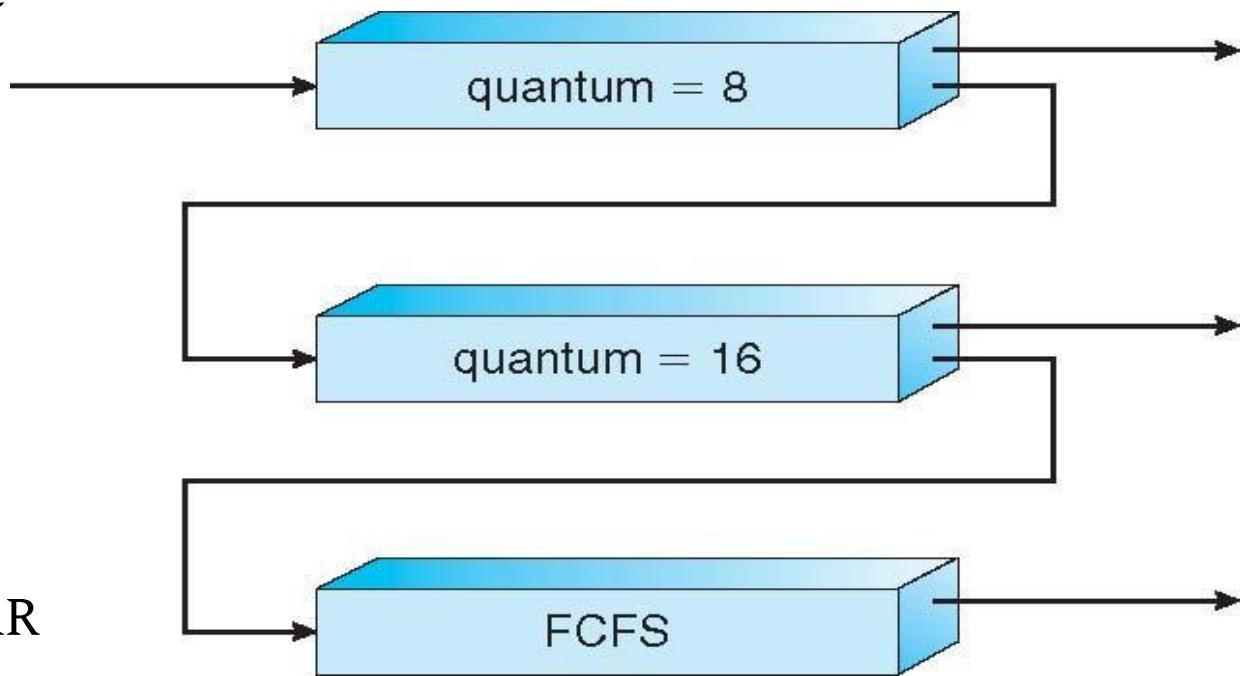
lowest priority

Multilevel Feedback Queue scheduling

- Multilevel Feedback Queue scheduling, allows a process to move between queues.
- If a process uses too much CPU time, it will be moved to a lower priority queue.
- Similarly, a process that waits too long in a lower-priority queue may me moved to a higher-priority queue.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

- Three queues: (can see the figure in next slide)
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served for RR
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 job is again served RR and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2

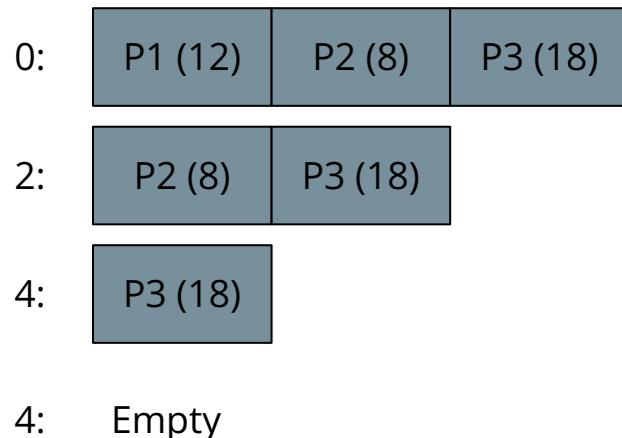


Example of Multilevel Feedback Queue

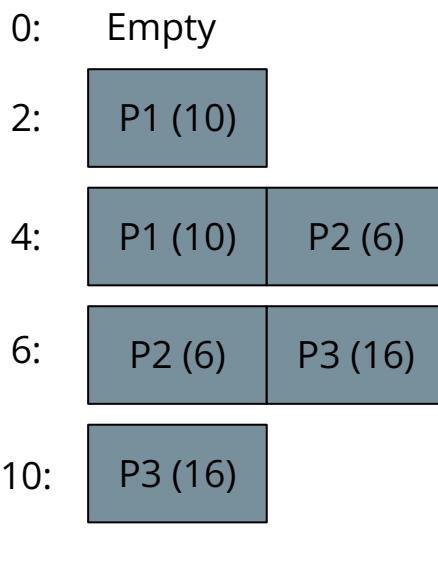
Problem 1:

Processes	Burst Time
P1	12
P2	8
P3	18

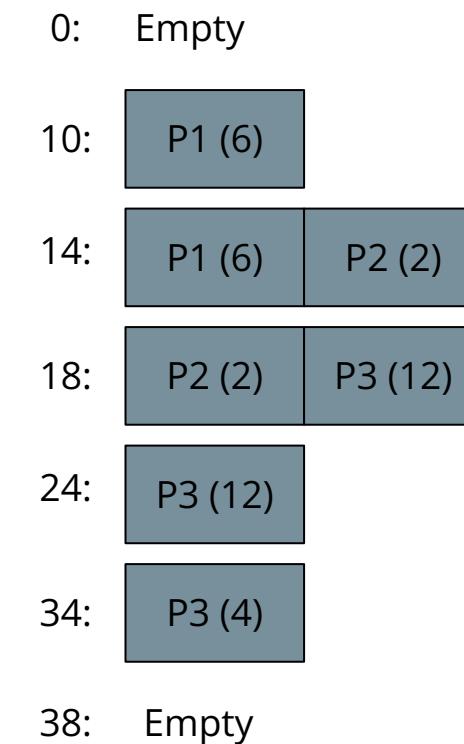
Queue 0 (Priority 0): Round-robin (quantum=2)



Queue 1 (Priority 1): Round-robin
(quantum=4)

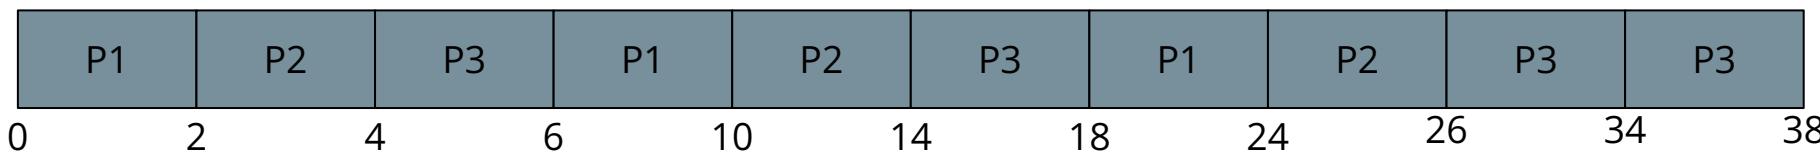


Queue 2 (Priority 2): Round-robin (quantum=8)



Gantt Chart:

CPU:



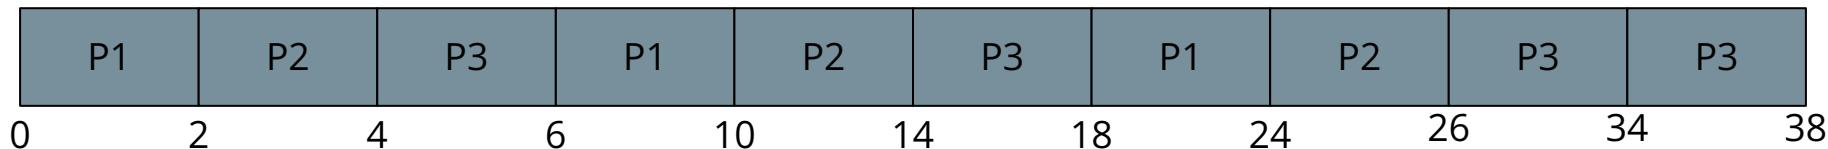
Example of Multilevel Feedback Queue

Problem 1:

Processes	Burst Time
P1	12
P2	8
P3	18

Gantt Chart:

CPU:



Waiting Time:

$$P1 = (0-0) + (6-2) + (18-10) = 12$$

$$P2 = (2-0) + (10-4) + (24-14) = 18$$

$$P3 = (4-0) + (14-6) + (26-18) + (34-34) = 20$$

$$Avg = (12+18+20)/3 = 16.67$$

Response Time:

$$P1 = (0-0) = 0$$

$$P2 = (2-0) = 2$$

$$P3 = (4-0) = 4$$

$$Avg = (0+2+4)/3 = 2$$

Turnaround Time:

$$P1 = (24-0) = 24$$

$$P2 = (26-0) = 26$$

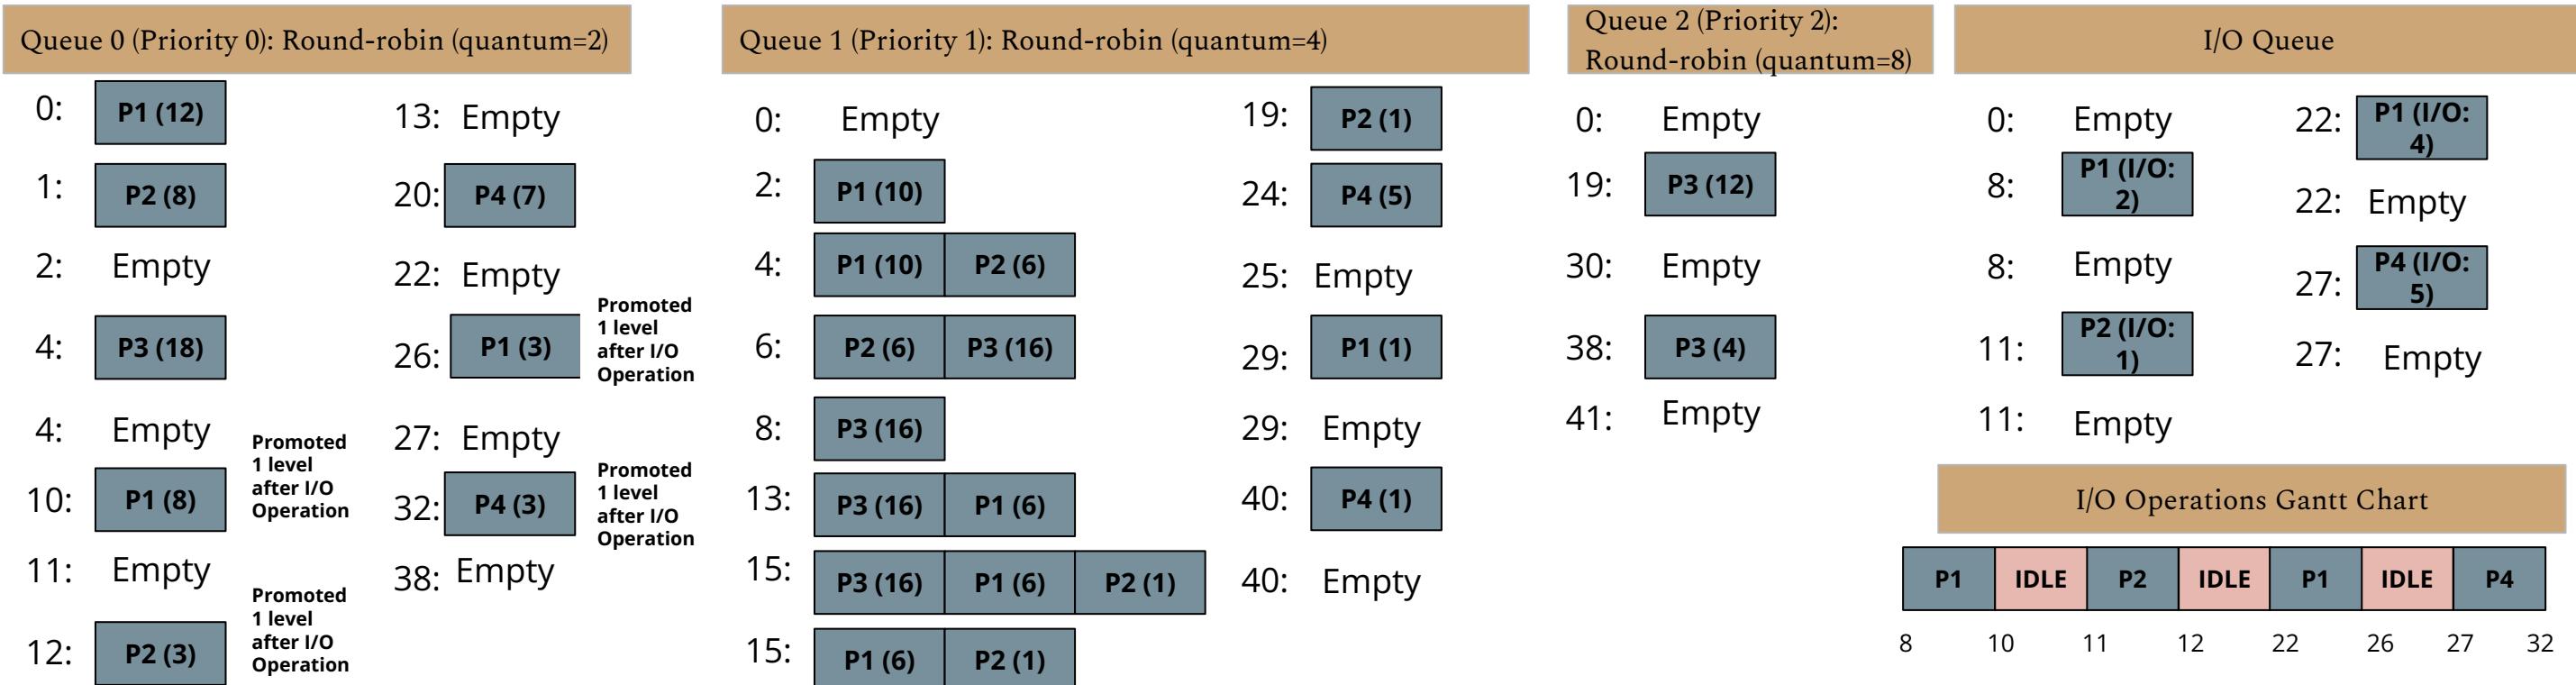
$$P3 = (38-0) = 38$$

$$Avg = (24+26+38)/3 = 29.33$$

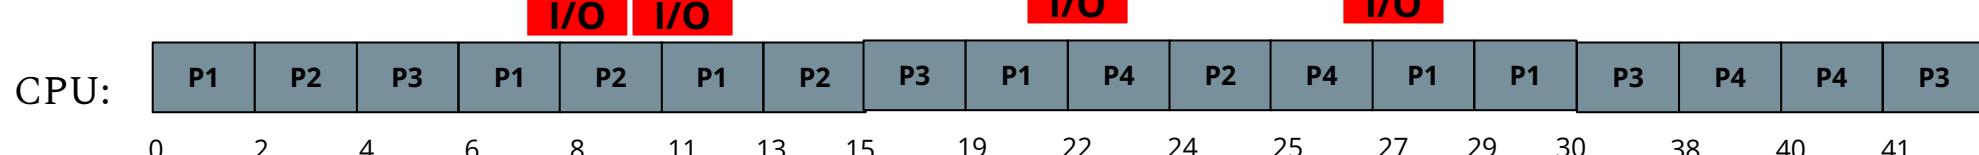
Example of Multilevel Feedback Queue

Problem 2:

Processes	Burst Time	Arrival Time	I/O Time
P1	12	0	6 [2s I/O operation after total 4s of CPU allocation & 4s I/O operation after total 9s of CPU allocation]
P2	8	1	1 [1s I/O operation after total 5s of CPU allocation]
P3	18	4	N/A
P4	7	20	5 [5s I/O operation after total 4s of CPU allocation]



Gantt Chart:

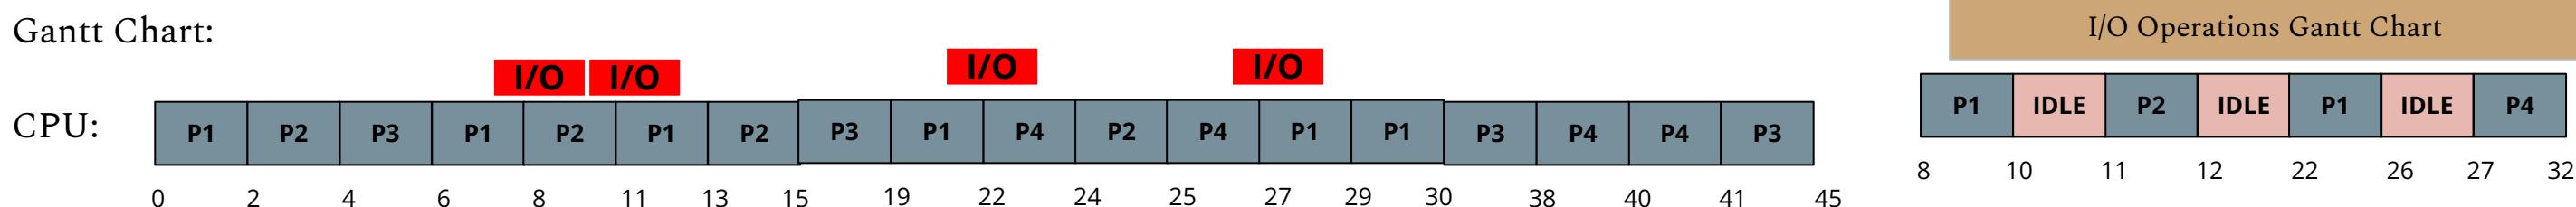


Example of Multilevel Feedback Queue

Problem 2:

Processes	Burst Time	Arrival Time	I/O Time
P1	12	0	6 [2s I/O operation after total 4s of CPU allocation & 4s I/O operation after total 9s of CPU allocation]
P2	8	1	1 [1s I/O operation after total 5s of CPU allocation]
P3	18	4	N/A
P4	7	20	5 [5s I/O operation after total 4s of CPU allocation]

Gantt Chart:



Waiting Time:

$$\begin{aligned} P1 &= (0-0)+(6-2)+(11-10)+(19-13)+(27-26)+(29-29) \\ &= 12 \end{aligned}$$

$$P2 = (2-1)+(8-4)+(13-12)+(24-15) = 15$$

$$P3 = (4-4)+(15-6)+(30-19)+(41-38) = 23$$

$$P4 = (22-20)+(25-24)+(38-32)+(40-40) = 9$$

$$Avg = (12+15+23+9)/4 = 14.75$$

Response Time:

$$P1 = (0-0) = 0$$

$$P2 = (2-1) = 1$$

$$P3 = (4-4) = 0$$

$$P4 = (22-20) = 2$$

$$Avg = (0+1+0+2)/4 = 0.75$$

Turnaround Time:

$$P1 = (30-0) = 30$$

$$P2 = (25-1) = 24$$

$$P3 = (45-4) = 41$$

$$P4 = (41-20) = 21$$

$$Avg = (30+24+41+21)/4 = 29$$